

Part 1 — AI Quiz Flow: Write-up

Competitive research

I audited 3 sponsored "free iq test" results and 1 top organic result (myiqtested.com) end-to-end as a user. None of the sponsored competitors used a traditional landing page — the first screen was already part of the quiz, and all three injected real-time percentile feedback mid-quiz to sustain engagement. All three asked for email right after the last question, before revealing the result, leveraging peak sunk-cost investment. One showed a price symbol before the quiz even started — an instant trust failure as a test user.

Funnel model & bottleneck

Transition	Assumed rate	Basis
Page visit → Quiz start	72%	Paid-search engagement, lifted by soft-entry design
Quiz start → Completion	65%	Published quiz-funnel benchmark (80M+ submissions)
Completion → Result shown	92%	Peak momentum; technical drop-off only
Result → Account created	42%	Single-field gate (vs. multi-field)
Modeled CVR	~18.1%	

Bottleneck ranking (impact of a +10pp improvement on overall CVR): Result→Account **+4.3pp** (highest leverage) > Quiz start→Complete +2.8pp > Visit→Quiz start +2.5pp. Every design decision below protects Result→Account specifically — and it's why that step is the first A/B test in Part 2.

Flow

Entry screen with the first quiz question visible (no separate landing/start step) → 8 quiz questions with mid-quiz percentile feedback → labor-illusion loading state → partial result (visual percentile curve, no number) with an inline email field → full result with qualitative category tiers → post-signup upsell teaser.

Key trade-offs

- Combined entry + first question, no landing page.** Removes the Landing→Start drop-off seen across all 3 competitors. Trade-off: less room for trust-building copy before the ask.
- Email gate is inline on the partial-result screen, not a separate screen/CTA.** Result→Account is the steepest modeled drop-off, so an unearned extra click there fights the funnel's own priority; result_viewed and account_created still fire as distinct events, so step-level measurement isn't lost.
- Percentile as a visual curve, category scores as qualitative tiers (Strong/Average/Needs practice), not raw numbers.** With 1-2 questions per category, a percentage implies false precision; the curve and tiers convey "where you stand" honestly instead.
- Zero monetary cues before account creation.** Informed by my own reaction, as a test user, to a competitor showing a price before the quiz started. Trade-off: no upsell surfaced pre-signup at all — deferred to a post-signup teaser screen instead.
- Fixed-height answer containers and a reserved toast slot.** No layout shift between questions, and the mid-quiz toast can never overlap question content, on any screen size — a mobile fix, since mobile is the weaker half of the funnel (~16.6% vs ~19.6% desktop CVR).

AI usage

I used Claude across three distinct roles, not a single "ask and accept" loop — each with a different level of oversight.

Research: Claude ran web searches for published quiz-funnel and paid-search benchmarks (Interact's 80M-submission dataset, Outgrow's single-field-gate findings) anchoring the funnel model above. I directed which claims needed sourcing and cross-checked the figures — the model separates benchmark-sourced numbers from my own estimated adjustments.

Building: Claude Code surfaced the file:// CORS constraint ruling out ES modules for a no-build flow (unprompted) and proposed the classic-script/namespace pattern used throughout src/. When implementing screens revealed an index-walk state machine couldn't express a mid-quiz email gate (needed for Part 2's Test 1), it proposed a resolver function — which I required it to verify against all 4 planned configs headlessly before accepting, not just demo the happy path.

Oversight: every non-trivial decision is logged, dated, with what was considered and why, in AI_WORKFLOW.md — including where I rejected its output. Clearest case: the initial partial-result screen used a blurred/locked exact score as the curiosity hook. I rejected this as a dark pattern (a fake blur reads as manipulative, not curious) and directed a percentile-curve visualization instead — real information honestly revealed, not a fake lock on data that doesn't yet exist.